



Simulation for model checking

Jeffrey Leek, Assistant Professor of Biostatistics
Johns Hopkins Bloomberg School of Public Health

Basic ideas

- Way back in the first week we talked about simulating data from distributions in R using the *rfoo* functions.
- In general simulations are way more flexible/useful
 - For bootstrapping as we saw in week 7
 - For evaluating models
 - For testing different hypotheses
 - For sensitivity analysis
- At minimum it is useful to simulate
 - A best case scenario
 - A few examples where you know your approach won't work
 - [The importance of simulating the extremes](#)

Simulating data from a model

Suppose that you have a regression model

$$Y_i = b_0 + b_1 X_i + e_i$$

Here is an example of generating data from this model where X_i and e_i are normal:



```
set.seed(44333)
x <- rnorm(50)
e <- rnorm(50)
b0 <- 1; b1 <- 2
y <- b0 + b1*x + e
```

Violating assumptions

```
set.seed(44333)
x <- rnorm(50)
e <- rnorm(50); e2 <- rcauchy(50)
b0 <- 1; b1 <- 2
y <- b0 + b1*x + e; y2 <- b0 + b1*x + e2
```

Violating assumptions

```
par(mfrow=c(1,2))  
plot(lm(y ~ x)$fitted,lm(y~x)$residuals,pch=19,xlab="fitted",ylab="residuals")  
plot(lm(y2 ~ x)$fitted,lm(y2~x)$residuals,pch=19,xlab="fitted",ylab="residuals")
```



Repeated simulations

```
set.seed(44333)
betaNorm <- betaCauch <- rep(NA,1000)
for(i in 1:1000){
  x <- rnorm(50); e <- rnorm(50); e2 <- rcauchy(50); b0 <- 1; b1 <- 2
  y <- b0 + b1*x + e; y2 <- b0 + b1*x + e2
  betaNorm[i] <- lm(y ~ x)$coeff[2]; betaCauch[i] <- lm(y2 ~ x)$coeff[2]
}
quantile(betaNorm)
```


0%	25%	50%	75%	100%
1.500	1.906	2.013	2.100	2.596

```
quantile(betaCauch)
```

0%	25%	50%	75%	100%
-278.352	1.130	1.965	2.804	272.391



Monte Carlo Error

```
boxplot(betaNorm,betaCauch,col="blue",ylim=c(-5,5))
```




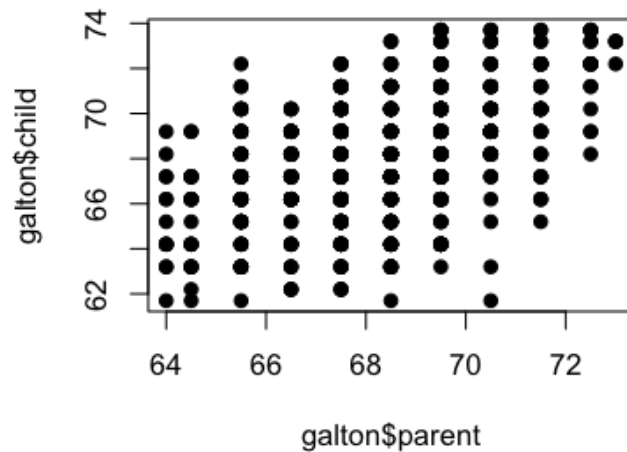
Simulation based on a data set

```
library(UsingR); data(galton); nobs <- dim(galton)[1]  
par(mfrow=c(1,2))  
hist(galton$child,col="blue",breaks=100)  
hist(galton$parent,col="blue",breaks=100)
```



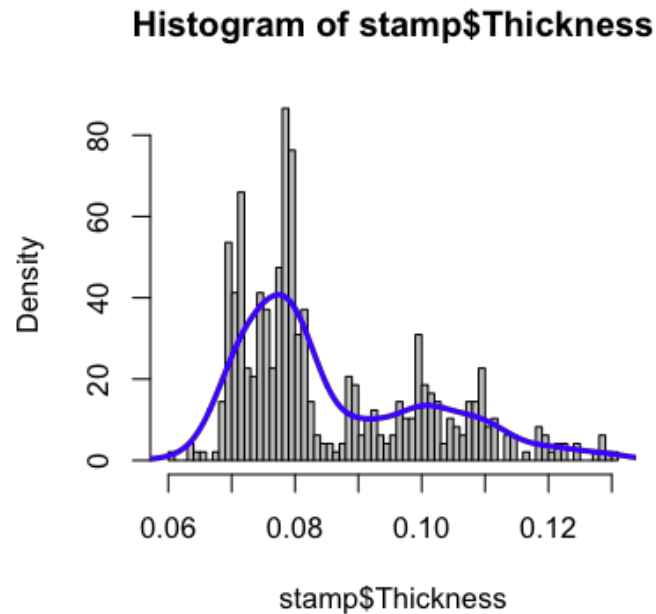
Calculating means,variances

```
lm1 <- lm(galton$child ~ galton$parent)
parent0 <- rnorm(nobs,sd=sd(galton$parent),mean=mean(galton$parent))
child0 <- lm1$coeff[1] + lm1$coeff[2]*parent0 + rnorm(nobs,sd=summary(lm1)$sigma)
par(mfrow=c(1,2))
plot(galton$parent,galton$child,pch=19)
plot(parent0,child0,pch=19,col="blue")
```



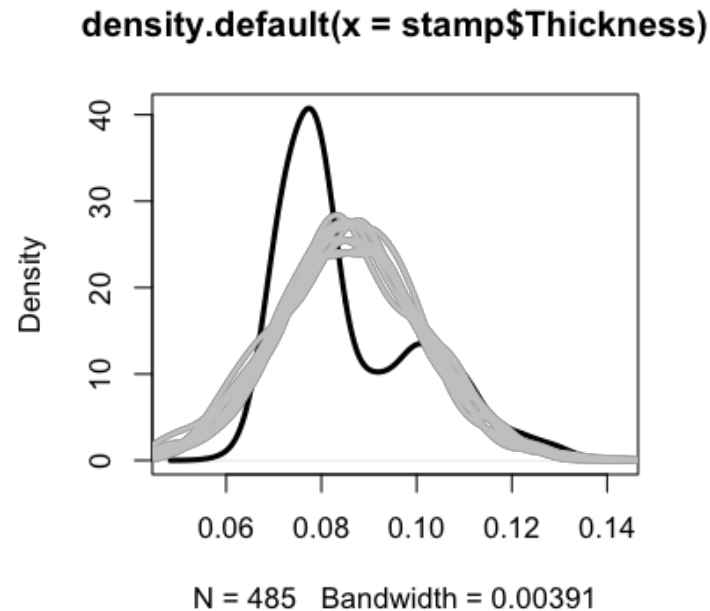
Simulating more complicated scenarios

```
library(bootstrap); data(stamp); nobs <- dim(stamp)[1]  
hist(stamp$Thickness,col="grey",breaks=100,freq=F)  
dens <- density(stamp$Thickness)  
lines(dens,col="blue",lwd=3)
```




A simulation that is too simple

```
plot(density(stamp$Thickness),col="black",lwd=3)
for(i in 1:10){
  newThick <- rnorm(nobs,mean=mean(stamp$Thickness),sd=sd(stamp$Thickness))
  lines(density(newThick),col="grey",lwd=3)
}
```




How density estimation works



http://en.wikipedia.org/wiki/File:Comparison_of_1D_histogram_and_KDE.png


Simulating from the density estimate

```
plot(density(stamp$Thickness),col="black",lwd=3)
for(i in 1:10){
  newThick <- rnorm(nobs,mean=stamp$Thickness,sd=dens$bw)
  lines(density(newThick),col="grey",lwd=3)
}
```



Increasing variability

```
plot(density(stamp$Thickness),col="black",lwd=3)
for(i in 1:10){
  newThick <- rnorm(nobs,mean=stamp$Thickness,sd=dens$bw*1.5)
  lines(density(newThick,bw=dens$bw),col="grey",lwd=3)
}
```



Notes and further resources

Notes

- Simulation can be applied to missing data problems - simulate what missing data might be
- Simulation values are often drawn from standard distributions, but this may not be appropriate
- Sensitivity analysis means trying different simulations with different assumptions and seeing how estimates change

Further resources

- [Advanced Data Analysis From An Elementary Point of View](#)
- [The design of simulation studies in medical statistics](#)
- [Simulation studies in statistics](#)